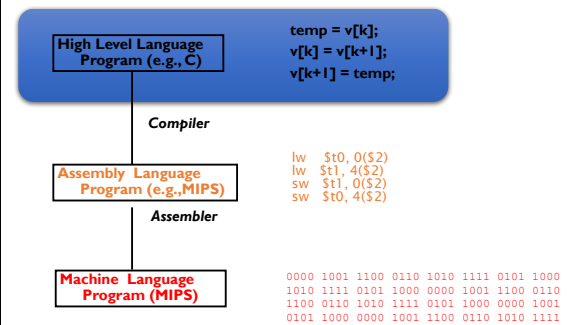


Lecture 2: Introduction to C

(CPEG323: Intro. to Computer System Engineering)

1

Levels of Program Code



The Lecture Plan

- You will not learn how to fully code in C in these lectures.
- We will only review a few key C concepts
 - Pointers
 - Arrays
 - Memory management

3

Compilation: Overview

- C *compilers* map C programs into architecture-specific machine code (string of 1s and 0s)
 - **compiling** .c files to .o files, then **linking** the .o files into executables; **Assembling** is also done (but is hidden, i.e., done automatically, by default)
- **Advantages**:
 - Excellent run-time performance
 - Allow us to exploit underlying features of the architecture
- **Disadvantages**:
 - Compiled files are **architecture-specific**, depending on CPU type and the operating system
 - Executable must be **rebuilt** on each new system
 - i.e., "porting your code" to a new architecture
 - "change → compile → run" iteration cycle can be slow, during the development cycle.

4

Actual C Code

```
#include <stdio.h>
#define REPEAT 5

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < REPEAT; i = i + 1) {
        printf("hello, world\n");
    }
    return 0;
}
```

Typed Variables in C

```
int    x = 2;
float  y = 1.618;
char   z = 'A';
```

Type	Description
int	signed integer
char	single text character or symbol
float	floating point non-integer number

- More about the integers:
 - The size of an integer is machine dependent! It is typically 4 bytes, but cannot assume it.
 - An integer can be either signed or unsigned (e.g., unsigned int).

6

C Syntax : Operators

- Arithmetic: +, -, *, /, %
- Assignment: =
- Bitwise
 - Logic: ~, &, |, ^
 - Shifts: <<, >>
- Boolean logic: !, &&, ||
- Equality testing: ==, !=
- Order relations: <, <=, >, >=
- ...

7

C Syntax : Control Flow

- if-else
 - if (expression) statement
 - if (expression) statement1 else statement2
- while
 - while (expression) statement
 - do statement while (expression);
- for
 - for (initialize; check; update) statement
- switch
 - switch (expression){
 - case const1: statements
 - case const2: statements
 - default: statements
 - break


8

Pointers

9

Address vs. Value


- Memory can be considered as a single huge array:
 - Each cell of the array has an address associated with it.
 - Each cell also stores some value.
- Don't confuse the **address** referring to a memory location with the **value** stored in that location.



10

Pointers

- A pointer is a variable that contains an address.
 - An *address* refers to a particular memory location, usually also associated with a variable name.



11

Pointer Operations

- `int *x;`
 - Declares variable x as the address of an integer.
- `x = &y;`
 - Assigns address of y to x.
- `z = *x;`
 - Assigns the value at address in x to z

12

Pointer Examples

```
int *p, x, y  p [?] x [?] y [?]
```

13

Example

```
int x=1, y=2, z[10];
int *ip;
```

```
ip = &x;
y = *ip;
*ip = 0;
ip = &z[0];
```

14

C pass parameters “by value”!

```
void AddOne (int x) {
    x = x + 1;
}
int y = 3;
AddOne(y);
```

Y remains equal to 3

Function AddOne gets a copy of the parameter, so changing the copy cannot change the original!

15

How to get a function to change a value?

- Pass a pointer!
 - Function accepts a pointer and then modifies value by dereferencing it.

```
void AddOne (int *p) {
    *p = *p + 1;
}
int y = 3;
AddOne(&y);
```

y is now equal to 4

16

Pointer Declaration and Allocation

- Declaring a pointer just allocates space to hold the pointer – it does not allocate the thing being pointed to!

- How to make it point to something meaningful?
 - Make it point to something that already exists

```
int *ptr, var1;
var1 = 5;
ptr = &var1;
```



- Allocate room in memory for something new that it will point to. (e.g., malloc)

17

An Example of buggy code

```
int *p;
*p = 5;
printf("%d\n", *p);
```

What is the result from executing this code?

- Prints 5
- Prints garbage
- Always crashes.
- Almost always crashes.

18

An Example of buggy code (Cont.)

```
int *p;
*p = 5;
printf("%d\n", *p);
```

What is the result from executing this code?

- Prints 5
- Prints garbage
- Always crashes.
- **Almost always crashes.**

19

Arrays

20

Basic Concepts of Arrays

- **Declaration:**
 - `int arr[2];`
 - `int arr[] = {795, 635};`
- **Accessing elements:**
 - `arr[k]`: returns the k^{th} element
 - Array size `n`: access entries 0 to `n-1`
 - Warning: An array in C does not know its own size, and its bounds are not checked! So, be careful with segmental faults and bus errors.

21

Pointer vs. Array

- Array variable is a “pointer” to the first (0^{th}) element
- `arr[i]` is treated as `*(arr+i)`
 - `arr[0]` is the same as `*arr`
 - `arr[2]` is the same as `*(arr+2)`
- Here are three equivalent ways to set all array elements to zero.
 - `for (i=0; i < size; i++) arr[i] = 0;`
 - `for (i=0; i < size; i++) *(arr+i) = 0;`
 - `for (p=a; p < arr + size; p++) *p = 0;`

22

Pointer Arithmetic

- A pointer is a memory address, so we can add to or subtract from it to move through in the memory space.

23

Pointer Arithmetic: An Example

Assume variable `a` is at address 100.

```
char *p;
char a;

p = &a;
```

```
int *p;
int a;

p = &a;
```

What does the following line yield?
`printf("%u %u\n", p, p+1);`

100 101

Adds `1*sizeof(char)` to the memory address

100 104

Adds `1*sizeof(int)` to the memory address

Pointer arithmetic should be used cautiously!!!

24


Pointer Arithmetic

- A pointer is a memory address, so we can add to or subtract from it to move through in the memory space.
- `p+1` means increments `p` by `sizeof(*p)`
 - i.e., moves pointer to the next array element.
- What is valid pointer arithmetic?
 - Add an integer to a pointer
 - Subtract 2 pointers (in the same array)
 - Compare pointers (`<`, `<=`, `==`, `!=`, `>`, `>=`)
 - Compare pointer to NULL (indicates that the pointer points to nothing)
- Everything else is illegal since it makes no sense:
 - Adding two pointers
 - Multiplying pointers
 - Subtract pointer from integer

25

Question

```
int main(void){
    int A[] = {5,10};
    int *p = A;
    printf("%u %d %d %d\n",p,*p,A[0],A[1]);
    p = p + 1;
    printf("%u %d %d %d\n",p,*p,A[0],A[1]);
    *p = *p + 1;
    printf("%u %d %d %d\n",p,*p,A[0],A[1]);
}
```



If the first printf outputs 100 5 5 10, what will the other two printf output?

1: 101 10 5 10 then 101 11 5 11
 2: 104 10 5 10 then 104 11 5 11
 3: 101 <other> 5 10 then 101 <3-others>
 4: 104 <other> 5 10 then 104 <3-others>
 5: One of the two printfs causes an ERROR
 6: I surrender!

26

Question (2)

How many of the following are **invalid**?

- pointer + integer
- integer + pointer
- pointer + pointer
- pointer - integer
- integer - pointer
- pointer - pointer
- compare pointer to pointer
- compare pointer to integer
- compare pointer to 0
- compare pointer to NULL

#invalid
1
2
3
4
5
6
7
8
9
10

27

Answer (2)

How many of the following are **invalid**?

- pointer + integer `ptr+1`
- integer + pointer `1+ptr`
- pointer + pointer `ptr+ptr`
- pointer - integer `ptr-1`
- integer - pointer `1-ptr`
- pointer - pointer `ptr-ptr`
- compare pointer to pointer `ptr1==ptr2`
- compare pointer to integer `ptr==1`
- compare pointer to 0 `ptr==0`
- compare pointer to NULL `ptr==NULL`

#invalid
1
2
3
4
5
6
7
8
9
(1)0

Summary

- All data is in memory
 - Each memory location has an address to use to refer to it and a value stored in it.
- Pointer is a variable whose value is an address
 - Operations:
 - * "follows" a pointer to its value
 - & gets the address of a value
 - Pointers and array variables are very similar.
 - Adding 1 to a pointer moves the pointer by the size of the thing it's pointing to.
 - When the pointers are useful?
 - If we want to pass a large array, it's easier to pass a pointer than the whole array.
 - In general, pointers allow cleaner, more compact code
 - So what are the drawbacks?
 - Pointers are probably the single largest source of bugs in C, so be careful anytime you deal with them, such as dangling references and memory leaks.


29

Memory Management

30

C Memory Layout

- There are four components in a program's address space.
 - stack**: local variables, grows downward
 - heap**: space requested for pointers via `malloc()`, grows upward
 - static data**: variables declared outside the main function, does not grow or shrink
 - code**: loaded when program starts, does not change.



31

Where are Variables Allocated?

- Declared outside a procedure
 - allocated in "static" storage
- Declared inside procedure
 - allocated on the "stack" and freed when procedure returns
 - `main()` is treated like a procedure
- Dynamically allocated via `malloc`:
 - Allocated on the "heap"

```
#include <stdio.h>


int myGlobal;

int main() {
    int myTemp;
    int *varDyn = malloc(sizeof(int));
}
```

32

Stack


- Stack consists of stack frames, i.e., a contiguous block of memory holding the local variables of a single procedure.
- A stack frame includes:
 - location of caller function (i.e., return address)
 - function arguments
 - local variables
- Stack pointer (SP) tells where the current stack frame is.
- When procedure ends, SP is moved back.



...

Stack – Last In, First Out

```
main ()
{ a(0);
  void a (int m)
  { b(1);
    }
  void b (int n)
  { c(2);
    }
  void c (int o)
  { d(3);
    }
  void d (int p)
  {
  }
```



34

What's Wrong with this Code?

```
int *getPtr() {
    int y;
    y = 3;
    return &y;
};


main () {
    int *stackAddr, content;
    stackAddr = getPtr();
    content = *stackAddr;
    printf("%d", content); /* 3 */
    content = *stackAddr;
    printf("%d", content); /*13451514 */
};
```

35

What's Wrong with this Code?

```
int *getPtr() {
    int y;
    y = 3;
    return &y;
};
```

```
main () {
    int *stackAddr, content;
    stackAddr = getPtr();
    content = *stackAddr;
    printf("%d", content); /* 3 */
    content = *stackAddr;
    printf("%d", content); /*13451514 */
};
```



Problem: *printf* overwrites stack frame

Never return pointers to local variable from functions
Your compiler will warn you about this – don't ignore such warnings!

36

Using the heap - Dynamic Memory Allocation

- **malloc(n):**
 - Allocate a block of uninitialized memory
 - Most often, malloc used to allocate space for an array of items.
 - `int *p = malloc (n*sizeof(int));`
 - Allocates space for n integers.
- **free(p):**
 - Releases memory allocated by malloc()
 - p is pointer containing the address originally returned by malloc()

37

Example

```
void foo() {
    int *p, *q, x, a[1];
    p = (int *) malloc (sizeof(int));
    q = &x;

    *p = 1;
    *q = 2;
    *a = 3;

    printf("*p:%u, p:%u, &p:%u\n", *p, p, &p);
    printf("*q:%u, q:%u, &q:%u\n", *q, q, &q);
    printf("*a:%u, a:%u, &a:%u\n", *a, a, &a);
}
```

... 12 16 20 24 28 32 36 40 44 48 52 56 60 64 68 ...

... [52] [32] [2] [3] ... unnamed-malloc-space

p q x a

*p:1, p:52, &p:24
 *q:2, q:32, &q:28
 *a:3, a:36, &a:36

38

Summary

- C has four pools of memory
 - Code
 - Static storage: global variable storage, permanent over entire program run
 - Stack: local variable storage, parameters, return address
 - Heap (dynamic storage): malloc() allocates space from here, free() returns it.
- Common Memory-related Bugs
 - Using uninitialized values
 - Accessing memory beyond your allocated region
 - Mismatched malloc/free pairs

39