

## Lecture 13: Pipeline Hazards

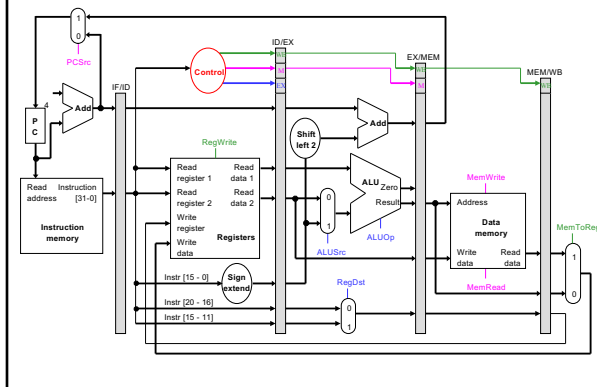
(CPEG323: Intro. to Computer System Engineering)

1

### Review

- Pipelining is a BIG idea
- Optimal Pipeline
  - Each instruction requires five stages, and five cycles, to complete.
  - Each stage uses different functional units of the datapath.
  - So we can execute up to five instructions in any clock cycle, with each instruction in a different stage and using different hardware.
- What makes this work well?
  - Similarities between instructions allow us to use same stages for all instructions (generally).
  - Each stage takes about the same amount of time as all others: little wasted time.

### The pipelined datapath




### Problems for Pipelining CPUs

- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
  - **Data hazards**: Instruction depends on result of prior instruction still in the pipeline
  - **Structural hazards**: HW cannot support some combination of instructions
  - **Control hazards**: Pipelining of branches causes later instruction fetches to wait for the result of the branch
- These might result in pipeline **stalls** or “**bubbles**” in the pipeline.

### Data Hazards

### Pipeline diagram review



- This diagram shows the execution of an ideal code fragment.
  - Each instruction needs a total of five cycles for execution.
  - One instruction begins on every clock cycle for the first five cycles.
  - One instruction completes on each cycle from that time on.

### Our examples are too simple

- Here is the example instruction sequence used to illustrate pipelining on the previous page.


```
lw $8, 4($29)
sub $2, $4, $5
and $9, $10, $11
or $16, $17, $18
add $13, $14, $0
```

- The instructions in this example are **independent**.
  - Each instruction reads and writes completely different registers.
  - Our datapath handles this sequence easily, as we saw last time.
- But most sequences of instructions are *not* independent!

### An example with dependencies


```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

### Data hazards in the pipeline diagram




- The SUB instruction does not write to register \$2 until clock cycle 5. This causes two **data hazards** in our current pipelined datapath.
  - The AND reads register \$2 in cycle 3. Since SUB hasn't modified the register yet, this will be the *old* value of \$2, not the new one.
  - Similarly, the OR instruction uses register \$2 in cycle 4, again before it's actually updated by SUB.

### Things that are okay



- The ADD instruction is okay, because of the register file design.
  - Registers are written at the beginning of a clock cycle.
  - The new value will be available by the end of that cycle.
- The SW is no problem at all, since it reads \$2 after the SUB finishes.


### Dependency arrows



- Arrows indicate the flow of data between instructions.
  - The tails of the arrows show when register \$2 is written.
  - The heads of the arrows show when \$2 is read.
- Any arrow that points backwards in time represents a **data hazard** in our basic pipelined datapath.

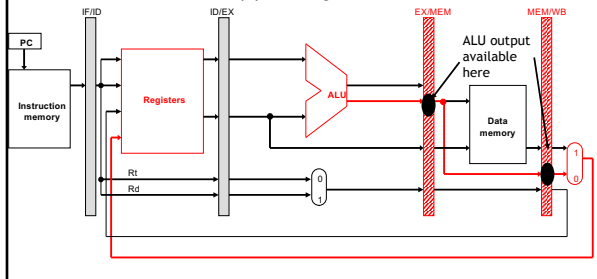
### Bypassing the register file

- The actual result \$1 - \$3 is computed in clock cycle 3, *before* it's needed in cycles 4 and 5.
- If we could somehow **bypass** the writeback and register read stages when needed, then we can eliminate these data hazards.
- Essentially, we need to pass the ALU output from SUB directly to the AND and OR instructions, without going through the register file.




### Pipeline Registers to the rescue!

- Pipeline stages communicate through pipeline registers: IF/ID ID/EX EX/MEM MEM/WB
- We “forward” data from pipeline registers to later instructions




### Forwarding

- The actual result \$1 - \$3 is computed in clock cycle 3, before it is needed in cycles 4 and 5
- We **forward** that value to later instructions, to prevent data hazards.
  - In clock cycle 4, AND gets the value \$1 - \$3 from the EX/MEM
  - In cycle 5, OR gets that same result from the MEM/WB

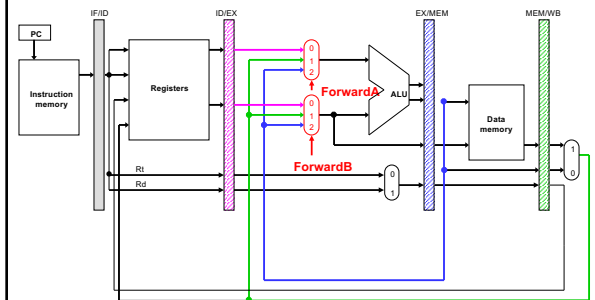


### Outline of forwarding hardware

- A **forwarding unit** selects the correct ALU inputs for the EX stage.
  - No hazard: ALU's operands come from the **register file**, like normal.
  - Data hazard: operands come from either the EX/MEM or MEM/WB pipeline registers instead.
- The ALU sources will be selected by two new multiplexers, with control signals named **ForwardA** and **ForwardB**.




### Simplified datapath with forwarding muxes



### Detecting EX/MEM Hazards

- In which stage cycle can we detect an impending hazard?



- Answer: cycle 3, when **sub** is in EX, **or** is in ID
  - Hazard because: ID/EX.rd == IF/ID.rs
- An **EX/MEM hazard** occurs between the instruction currently in its EX stage and the previous instruction if:
  - The previous instruction will write to the register file, *and*
  - The destination is one of the ALU source registers in the EX stage

17


### EX/MEM data hazard equations

- The first ALU source comes from the pipeline register when necessary.

if (EX/MEM.RegWrite = 1  
and EX/MEM.RegisterRd = ID/EX.RegisterRs)  
then ForwardA = 2

- The second ALU source is similar.

if (EX/MEM.RegWrite = 1  
and EX/MEM.RegisterRd = ID/EX.RegisterRt)  
then ForwardB = 2




### Detecting MEM/WB data hazards

- A **MEM/WB hazard** may occur between an instruction in the EX stage and the instruction from *two* cycles ago.
- One new problem is if a register is updated twice in a row.

```
add $1, $2, $3
add $1, $1, $4
sub $5, $5, $1
```

- Register \$1 is written by *both* of the previous instructions; from which instruction should it receive its value?




### Detecting MEM/WB data hazards

- A **MEM/WB hazard** may occur between an instruction in the EX stage and the instruction from *two* cycles ago.
- One new problem is if a register is updated twice in a row.

```
add $1, $2, $3
add $1, $1, $4
sub $5, $5, $1
```

- Register \$1 is written by *both* of the previous instructions, but only the most recent result (from the second ADD) should be forwarded.



### MEM/WB hazard equations

- Here is an equation for detecting and handling MEM/WB hazards for the first ALU source.

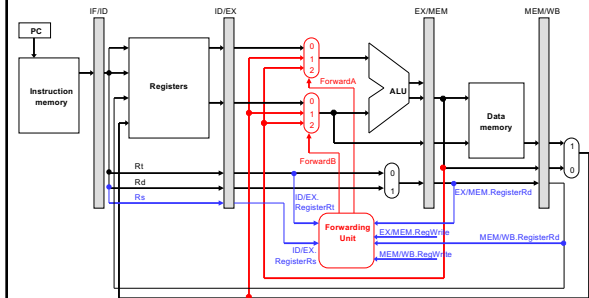
```
if (MEM/WB.RegWrite = 1
    and MEM/WB.RegisterRd = ID/EX.RegisterRs
    and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRs or EX/MEM.RegWrite = 0))
    then ForwardA = 1
```

- The second ALU operand is handled similarly.

```
if (MEM/WB.RegWrite = 1
    and MEM/WB.RegisterRd = ID/EX.RegisterRt
    and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRt or EX/MEM.RegWrite = 0))
    then ForwardB = 1
```

- Handled by a forwarding unit which uses the control signals stored in pipeline registers to set the values of ForwardA and ForwardB

### Simplified datapath with forwarding



### The forwarding unit

- The forwarding unit has several control signals as inputs.

ID/EX.RegisterRs	EX/MEM.RegisterRd	MEM/WB.RegisterRd
ID/EX.RegisterRt	EX/MEM.RegWrite	MEM/WB.RegWrite

(The two RegWrite signals are not shown in the diagram, but they come from the control unit.)

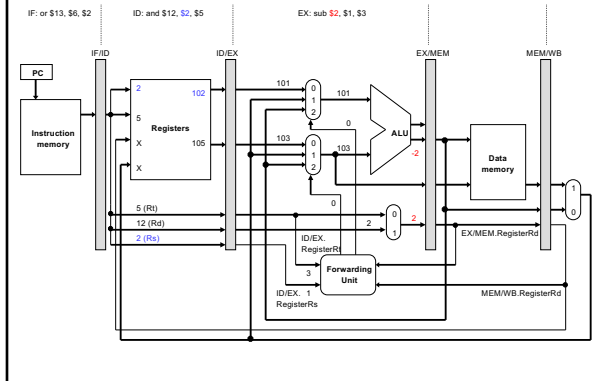
- The forwarding unit outputs are selectors for the **ForwardA** and **ForwardB** multiplexers attached to the ALU. These outputs are generated from the inputs using the equations on the previous pages.
- Some new buses route data from pipeline registers to the new muxes.

### Example

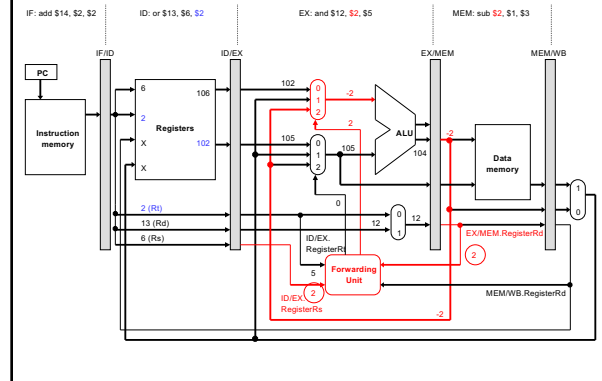
```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

- Assume again each register initially contains its number plus 100.
  - After the first instruction, \$2 should contain -2 (101 - 103).
  - The other instructions should all use -2 as one of their operands.
- We'll try to keep the example short.
  - Assume no forwarding is needed except for register \$2.
  - We'll skip the first two cycles, since they're the same as before.

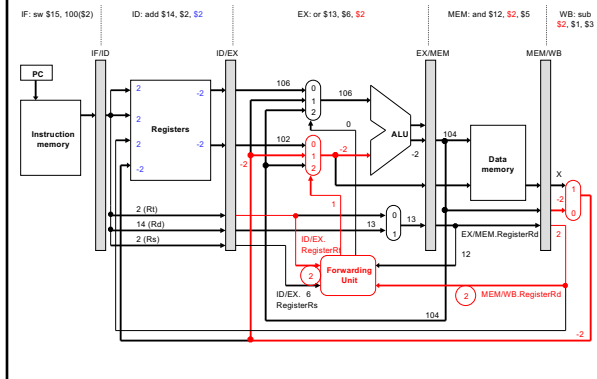
### Clock cycle 3



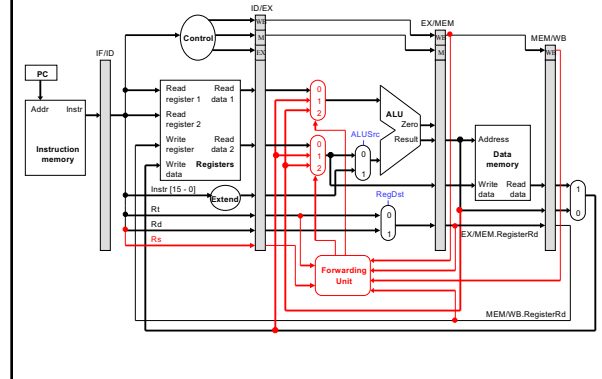
### Clock cycle 4: forwarding \$2 from EX/MEM



### Clock cycle 5: forwarding \$2 from MEM/WB




### Complete pipelined datapath...so far

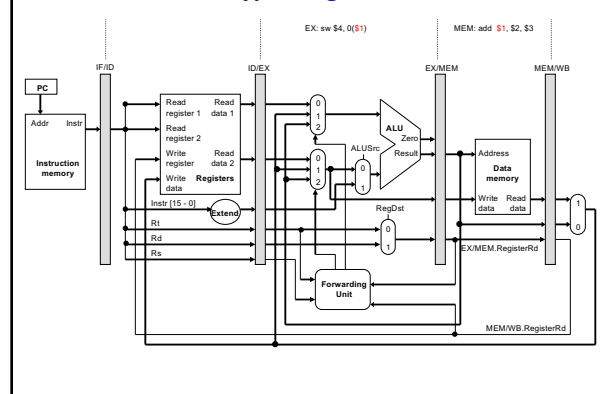


### What about stores?

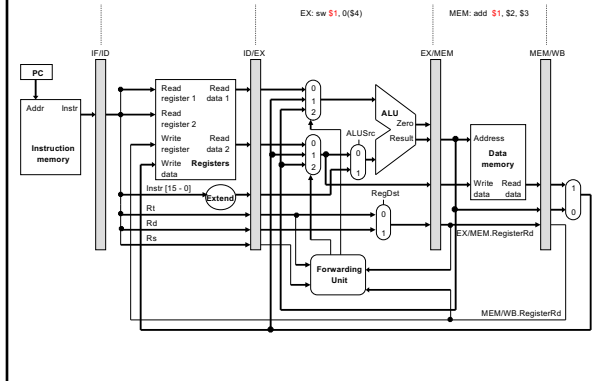
- Two "easy" cases:



### Store Bypassing: Version 1




### Store Bypassing: Version 2



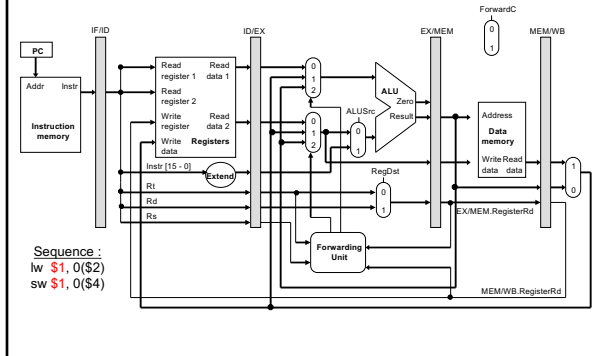
### What about stores?

- A harder case:




- In what cycle is the load value available?
  - End of cycle 4
- In what cycle is the store value needed?
  - Start of cycle 5
- What do we have to add to the datapath?

### Load/Store Bypassing: Extend the Datapath



### What about loads?


- Consider the instruction sequence shown below:
  - The load data doesn't come from memory until the *end* of cycle 4.
  - But the AND needs that value at the *beginning* of the same cycle!
- This is a "true" data hazard—the data is not available when we need it.



- We call this a load-use hazard.

### Stalling


- The easiest solution is to **stall** the pipeline.
- We could delay the AND instruction by introducing a one-cycle delay into the pipeline, sometimes called a **bubble**.



- Notice that we're still using forwarding in cycle 5, to get data from the MEM/WB pipeline register to the ALU.

### Stalling and forwarding


- Without forwarding, we'd have to stall for *two* cycles to wait for the LW instruction's writeback stage.



- In general, you can always stall to avoid hazards—but dependencies are very common in real code, and stalling often can reduce performance by a significant amount.

### Stalling delays the entire pipeline

- If we delay the second instruction, we'll have to delay the third one too.
  - This is necessary to make forwarding work between AND and OR.
  - It also prevents problems such as two instructions trying to write to the same register in the same cycle.



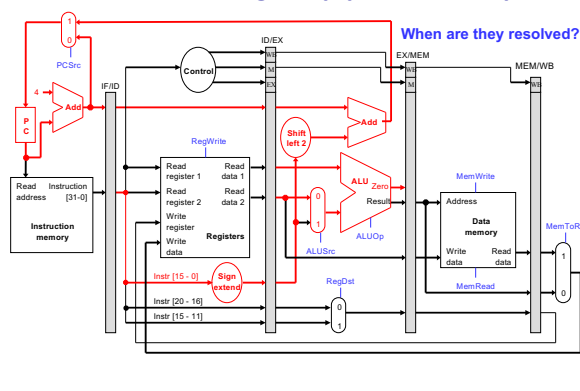
### Summary

- In real code, most instructions are dependent upon other ones.
  - This can lead to **data hazards** in our original pipelined datapath.
  - Instructions can't write back to the register file soon enough for the next two instructions to read.
- Forwarding** eliminates data hazards involving arithmetic instructions.
  - The forwarding unit detects hazards by comparing the destination registers of previous instructions to the source registers of the current instruction.
  - Hazards are avoided by grabbing results from the pipeline registers *before* they are written back to the register file.
- Stalling**
  - Forwarding can't save us in some cases involving `lw`.

### Structural Hazards


### Control Hazards

### Branches in the original pipelined datapath




### Branches

- Most of the work for a branch computation is done in the EX stage.
  - The branch target address is computed.
  - The source registers are compared by the ALU, and the Zero flag is set or cleared accordingly.
- Thus, the branch decision cannot be made until the end of the EX stage.
  - But we need to know which instruction to fetch next, in order to keep the pipeline running!
  - This leads to what's called a **control hazard**.



### Stalling is one solution


- Again, stalling is always one possible solution.



- Here we just stall until cycle 4, after we do make the branch decision.


### Branch prediction

- Another approach is to *guess* whether or not the branch is taken.
  - In terms of hardware, it's easier to assume the branch is *not* taken.
  - This way we just increment the PC and continue execution, as for normal instructions.
- If we're correct, then there is no problem and the pipeline keeps going at full speed.



### Branch misprediction

- If our guess is wrong, then we would have already started executing two instructions incorrectly. We'll have to discard, or **flush**, those instructions and begin executing the right ones from the branch target address, Label.



### Performance gains and losses

- Overall, branch prediction is worth it.
  - Mispredicting a branch means that two clock cycles are wasted.
  - But if our predictions are even just occasionally correct, then this is preferable to stalling and wasting two cycles for *every* branch.
- All modern CPUs use branch prediction.
  - Accurate predictions are important for optimal performance.
  - Most CPUs predict branches dynamically—statistics are kept at run-time to determine the likelihood of a branch being taken.
- The pipeline structure also has a big impact on branch prediction.
  - A longer pipeline may require more instructions to be flushed for a misprediction, resulting in more wasted time and lower performance.
  - We must also be careful that instructions do not modify registers or memory before they get flushed.

### Summary

- Three kinds of hazards conspire to make pipelining difficult.
- Structural hazards** result from not having enough hardware available to execute multiple instructions simultaneously.
  - These are avoided by adding more functional units (e.g., more adders or memories) or by redesigning the pipeline stages.
- Data hazards** can occur when instructions need to access registers that haven't been updated yet.
  - Hazards from R-type instructions can be avoided with forwarding.
  - Loads can result in a "true" hazard, which must stall the pipeline.
- Control hazards** arise when the CPU cannot determine which instruction to fetch next.
  - We can minimize delays by doing branch tests earlier in the pipeline.
  - We can also take a chance and predict the branch direction, to make the most of a bad situation.
- Number of cycles for  $N$  instructions on a  $k$  stage pipeline:
 
$$N + k - 1 + s + f$$
 where  $s = \text{\#stalls}$ ,  $f = \text{\#flushes}$

### Reading

- 5<sup>th</sup> Edition: 4.7,4.8