

CPEG 422/622 EMBEDDED SYSTEMS DESIGN

MW 3:35 – 4:50, SHARP 116

Chengmo Yang

chengmo@udel.edu

Evans 201C

1

LECTURE 14

HW/SW CO-DESIGN

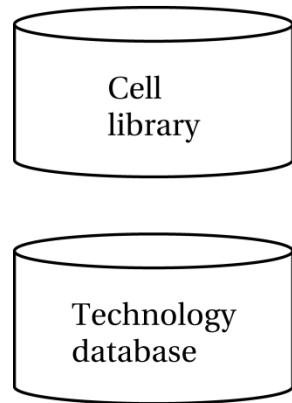
2

OUTLINE

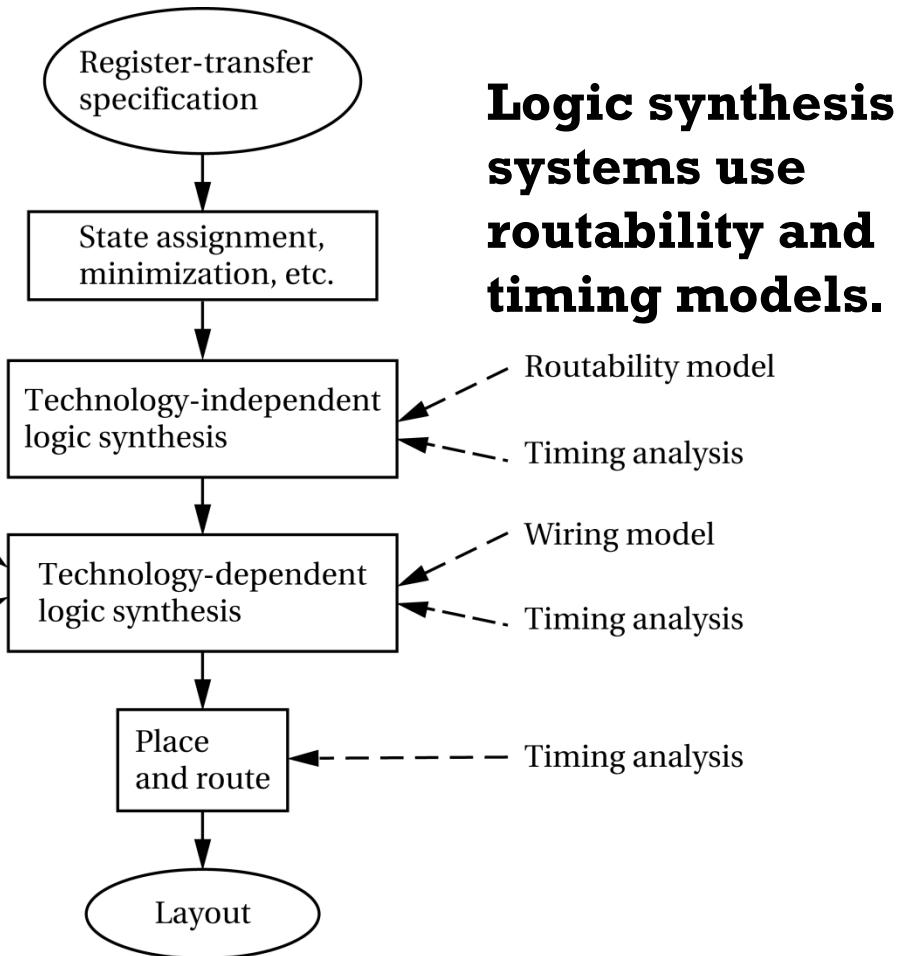
- Various design flows
 - Hardware design flow
 - Platform-based design
 - HW/SW co-design flow
- Partitioning
- Scheduling

HARDWARE DESIGN FLOW

Cell libraries describe the cells used to compose designs.



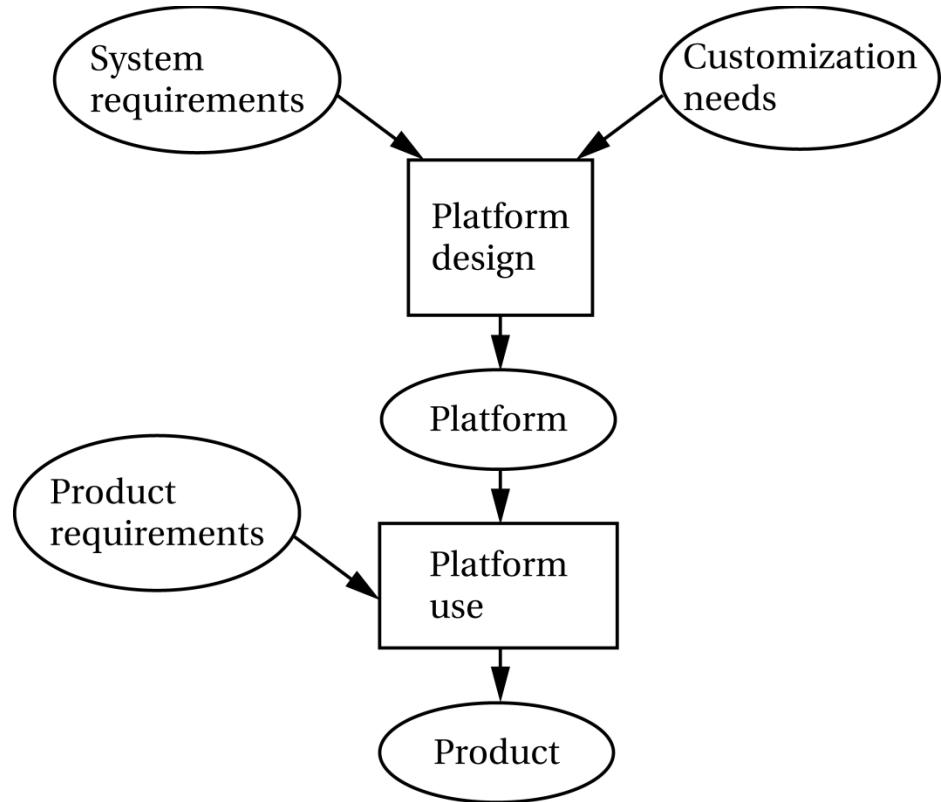
Technology databases capture manufacturing process information.



Logic synthesis systems use routability and timing models.

PLATFORM-BASED DESIGN

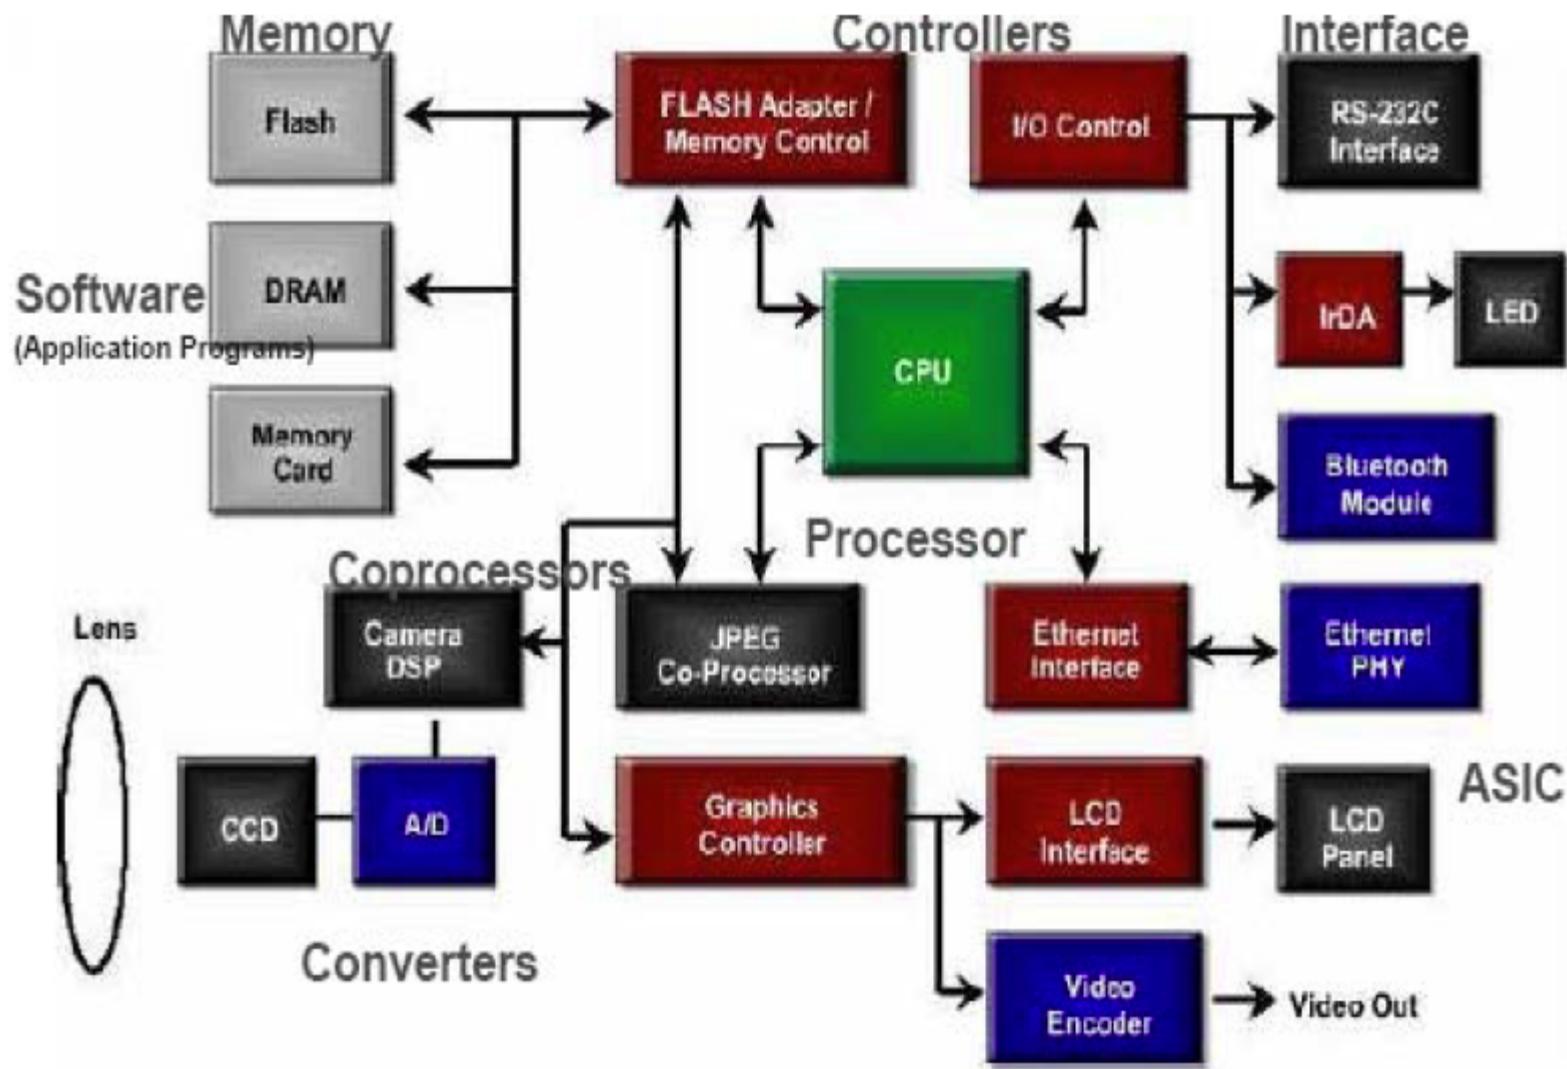
- Platform includes hardware, supporting software.
- Two stage process:
 - Design the platform.
 - Use the platform.
- Platform can be reused to host many different systems.



DESIGN PLATFORMS

- Different levels of integration:
 - PC + board.
 - Custom board with CPU + FPGA or ASIC.
 - Platform FPGA.
 - System-on-chip (SoC).

HARDWARE PLATFORM ARCHITECTURE

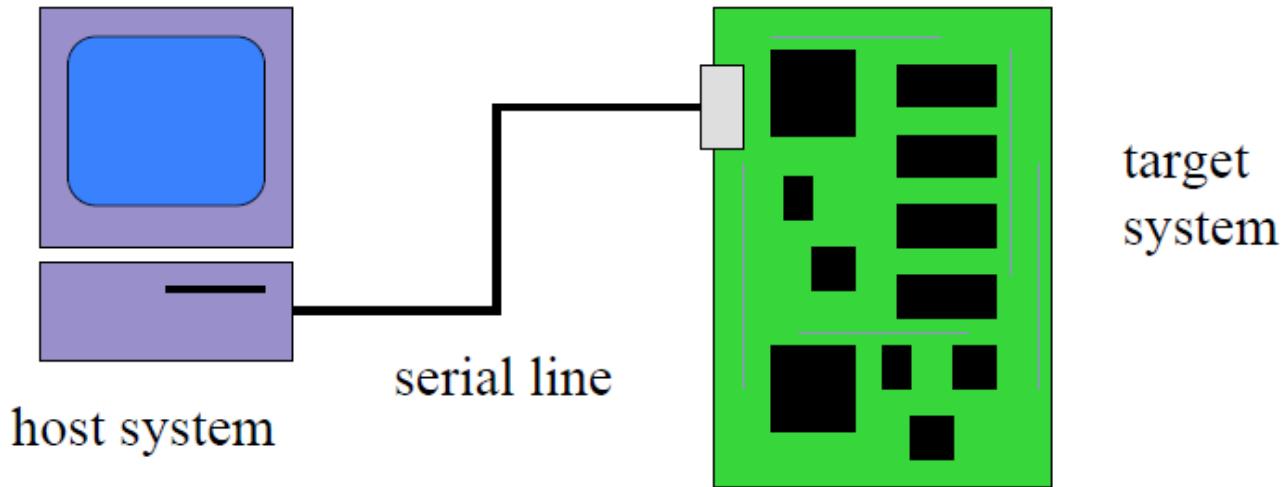


THE PC AS A PLATFORM

- Advantages:
 - Cheap and easy to get
 - Rich and familiar software environment
- Disadvantages:
 - Requires a lot of hardware resources
 - Not well-adapted to real-time

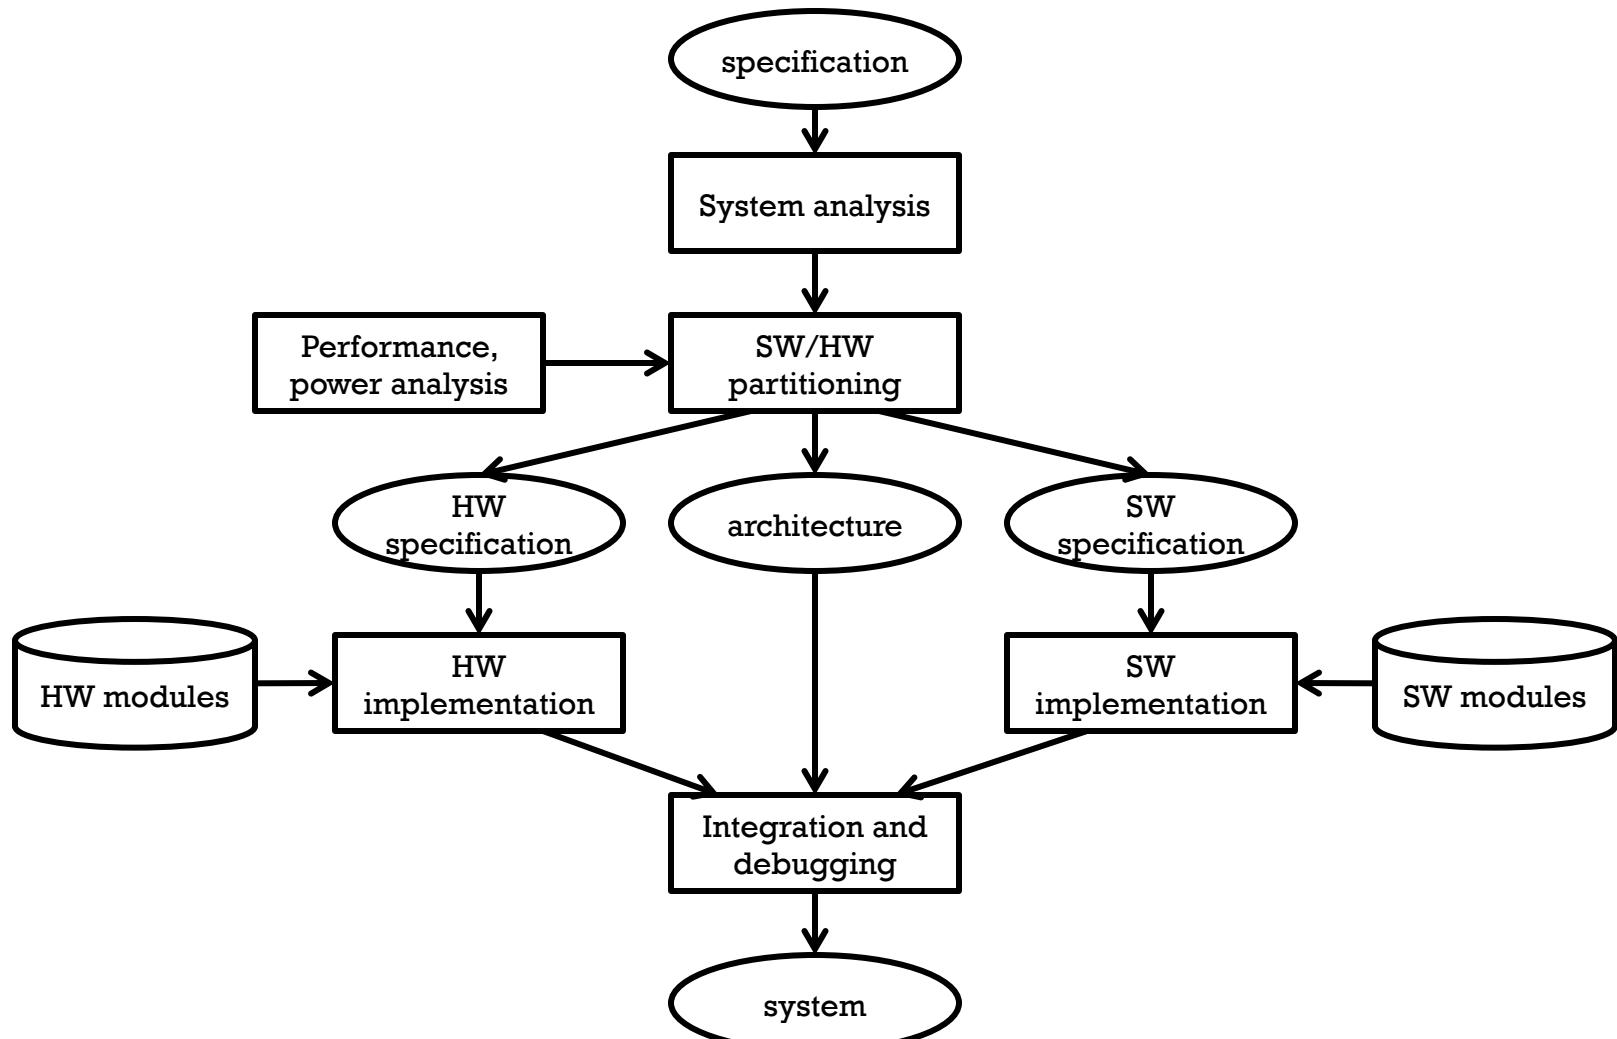
HOST / TARGET DESIGN

- Use a host system to prepare software for target system

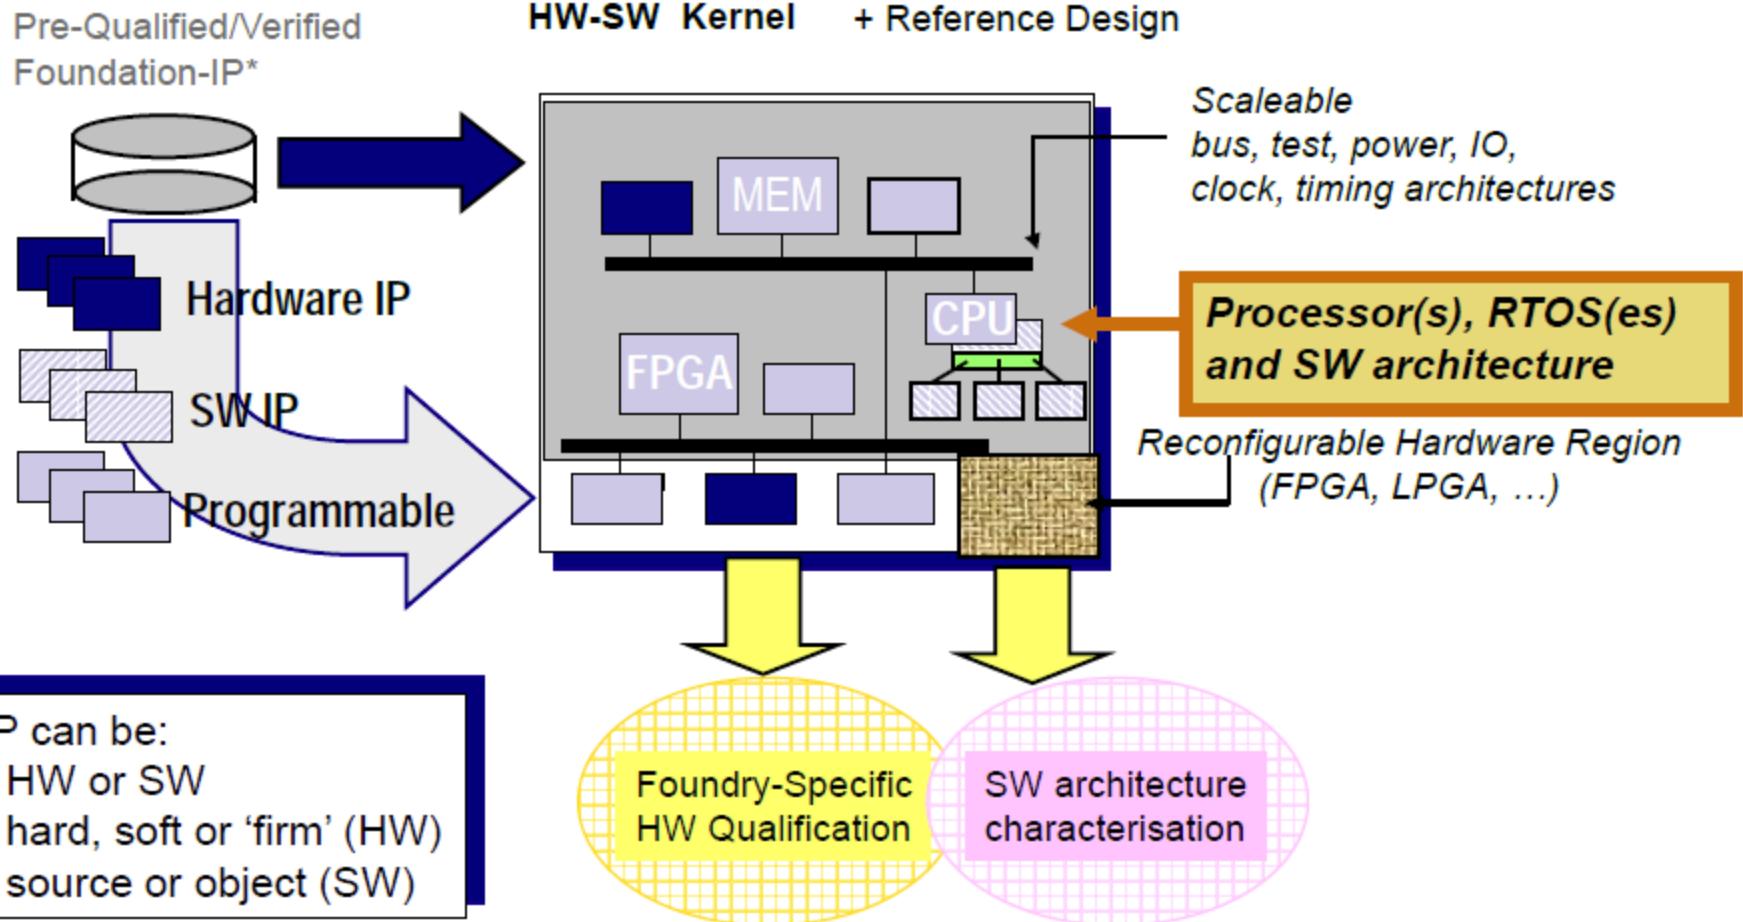


- Host-Based Tools
 - Cross compiler: Compiles code on host for target system
 - Cross debugger: Displays target state, allows target system to be controlled

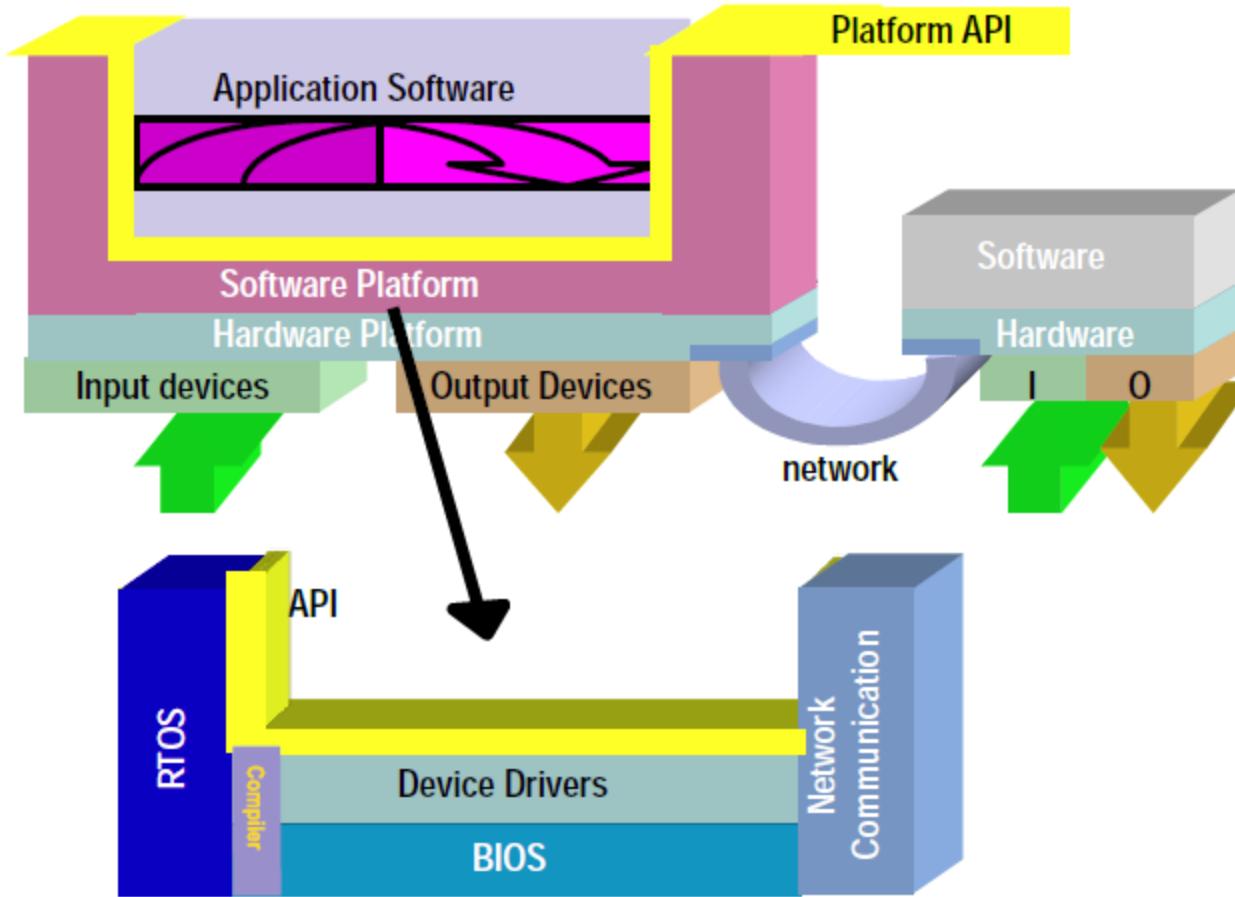
HARDWARE/SOFTWARE CO-DESIGN FLOW



HW-CENTRIC VIEW OF A PLATFORM



SW-CENTRIC VIEW OF PLATFORMS



HW/SW CODESIGN ISSUES

- **Task level concurrency management**

- Which tasks in the final system?

- **High level transformations**

- Transformation outside the scope of traditional compilers

- **Hardware/software partitioning**

- Which operation mapped to hardware, which to software?

- **Compilation**

- Hardware-aware compilation

- **Scheduling**

- Performed several times, with varying precision

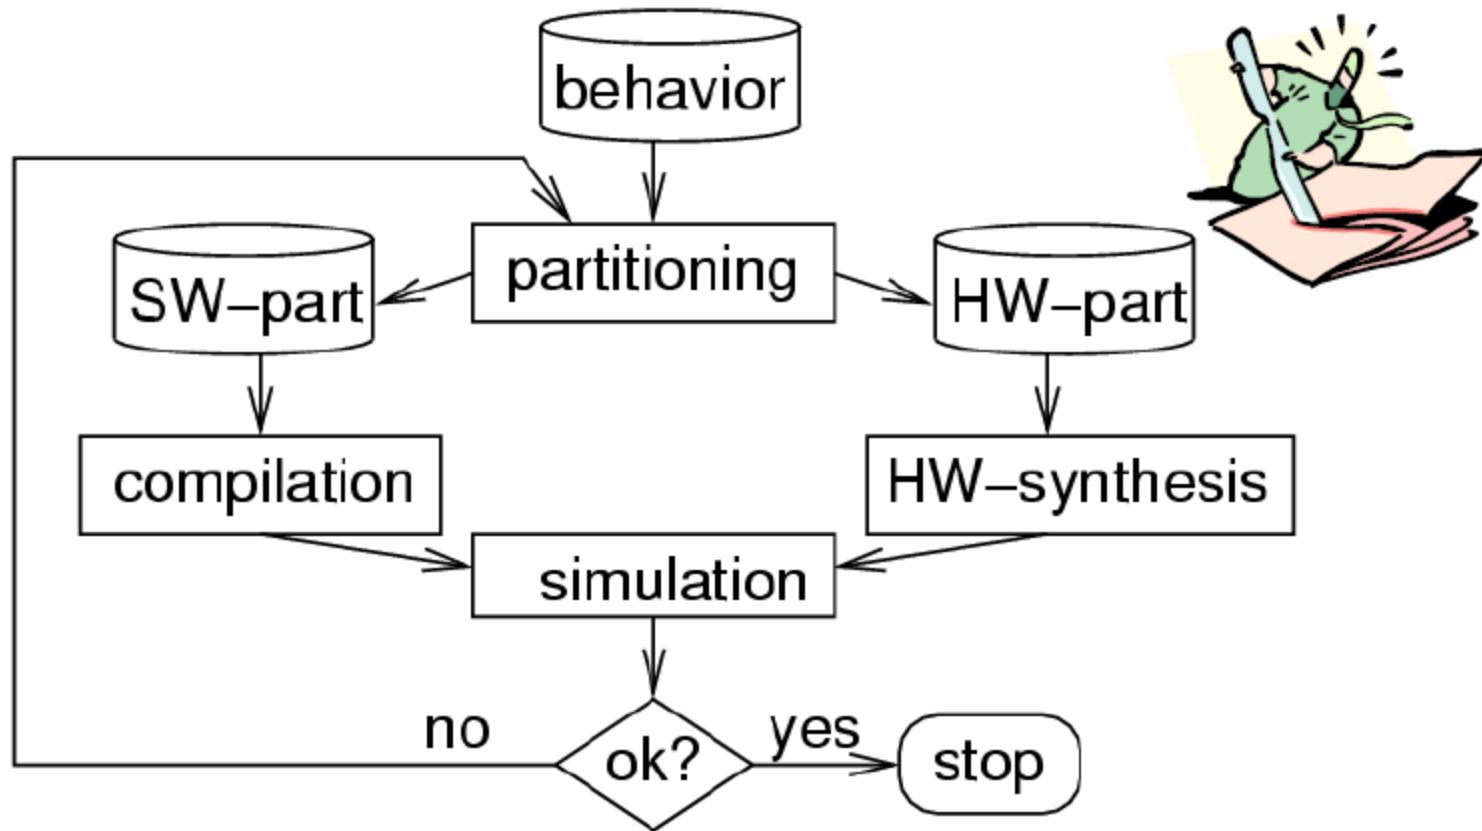
- **Design space exploration**

- Set of possible designs, not just one.

PARTITIONING

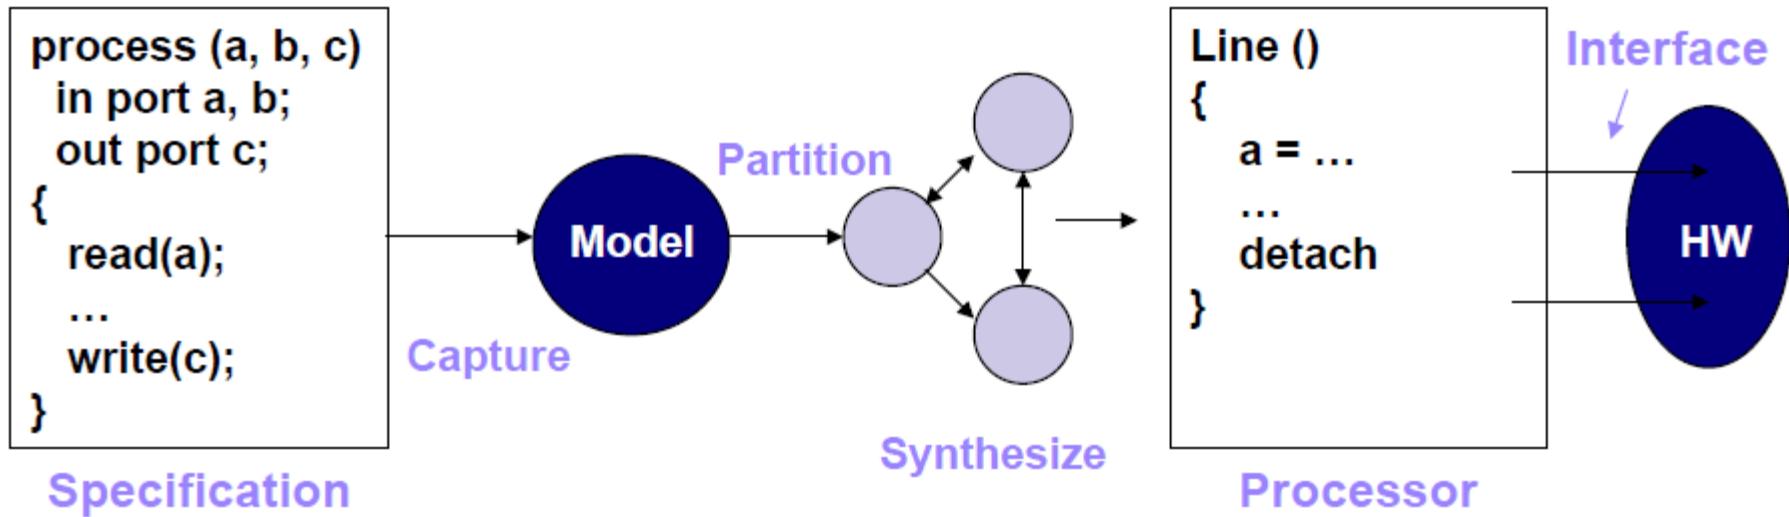
14

SOFTWARE OR HARDWARE?



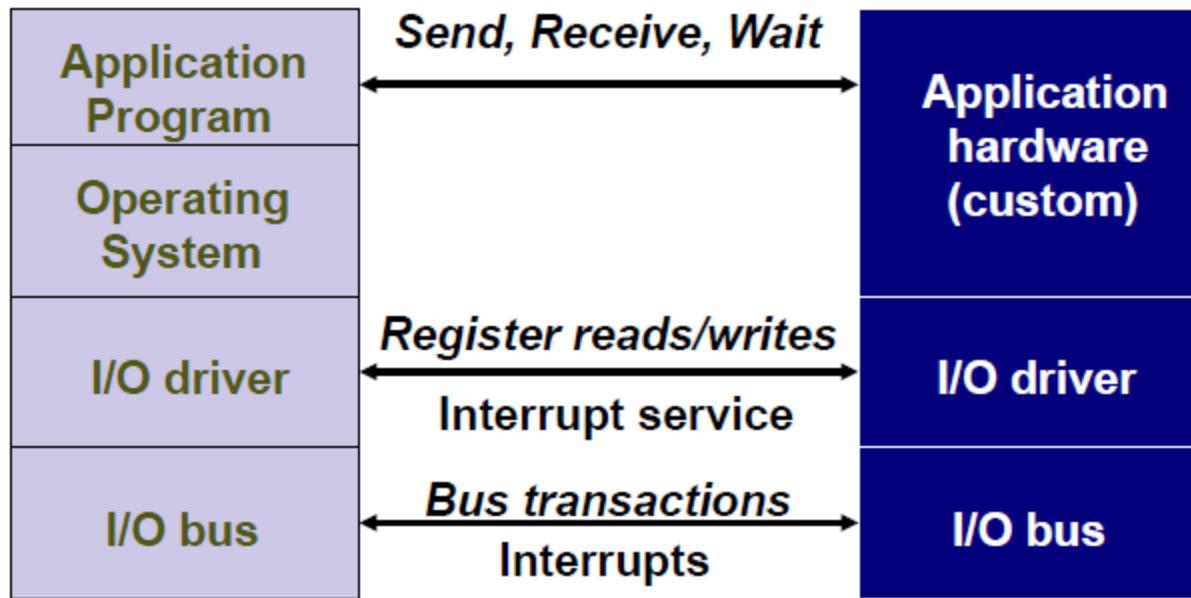
Decision based on hardware/ software partitioning

SYSTEM PARTITIONING



- Good partitioning mechanism:
 - Minimize communication across bus
 - Allows parallelism -> both HW & CPU operating concurrently
 - Near peak processor utilization at all times

DETERMINING COMMUNICATION LEVEL



- Easier to program at application level
 - (send, receive, wait) but difficult to predict
- More difficult to specify at low level
 - Difficult to extract from program but timing and resources easier to predict

PARTITIONING COSTS

- Software Resources

- Performance and power consumption
- Lines of code –development and testing cost
- Cost of components

- Hardware Resources

- Fixed number of gates, limited memory & I/O
- Difficult to estimate timing for custom hardware
- Recent design shift towards IP
 - Well-defined resource and timing characteristics

EXAMPLE: HW & SW FOUNDRIES

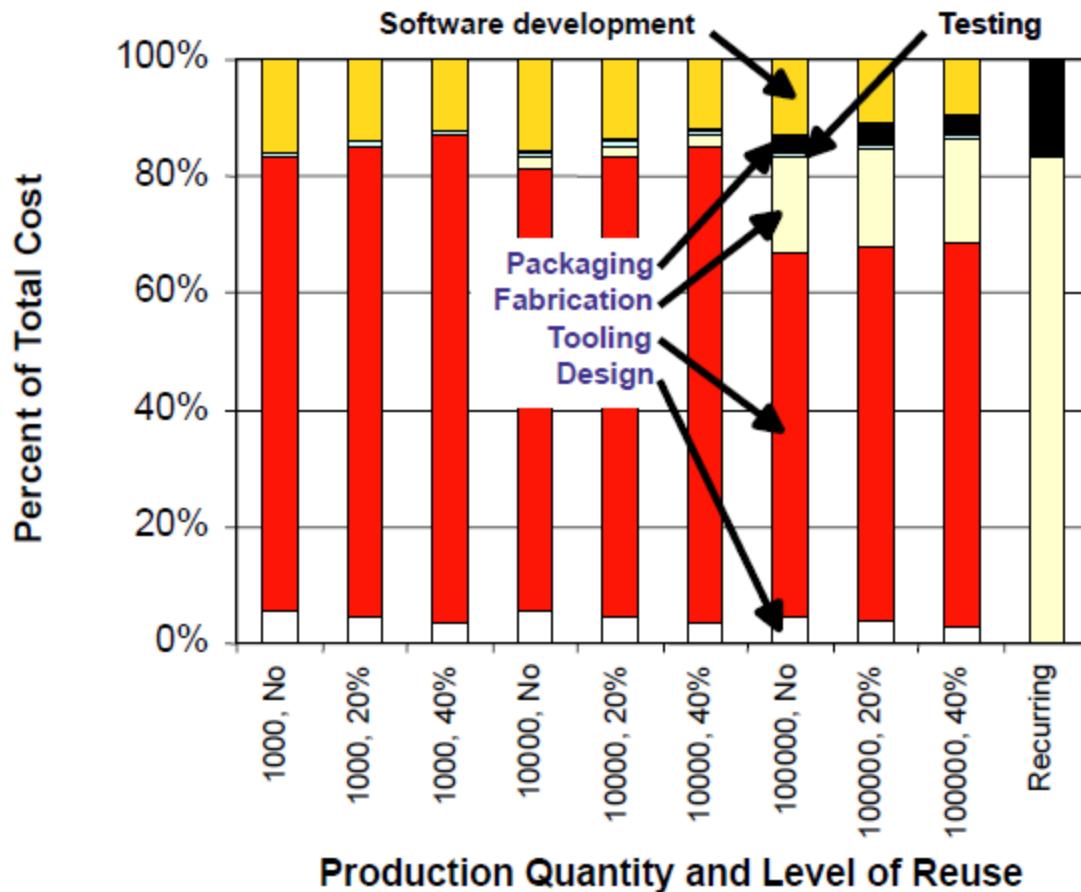
HW foundries

- HW1
 - LSI Logic ASIC Wafer Foundry Data
 - 0.18 μm feature size
 - 8 inch wafers
 - 6 layers
 - TSMC 018 Wafer Processing
- HW2
 - Samsung Semiconductor ASIC Wafer Foundry Data
 - 0.35 μm feature size
 - 6 inch wafers
 - 4 layers
 - TSMC 035 Wafer Processing

SW foundries

- SW1
 - Nominal to High development effort
- SW2
 - Low to Nominal development effort

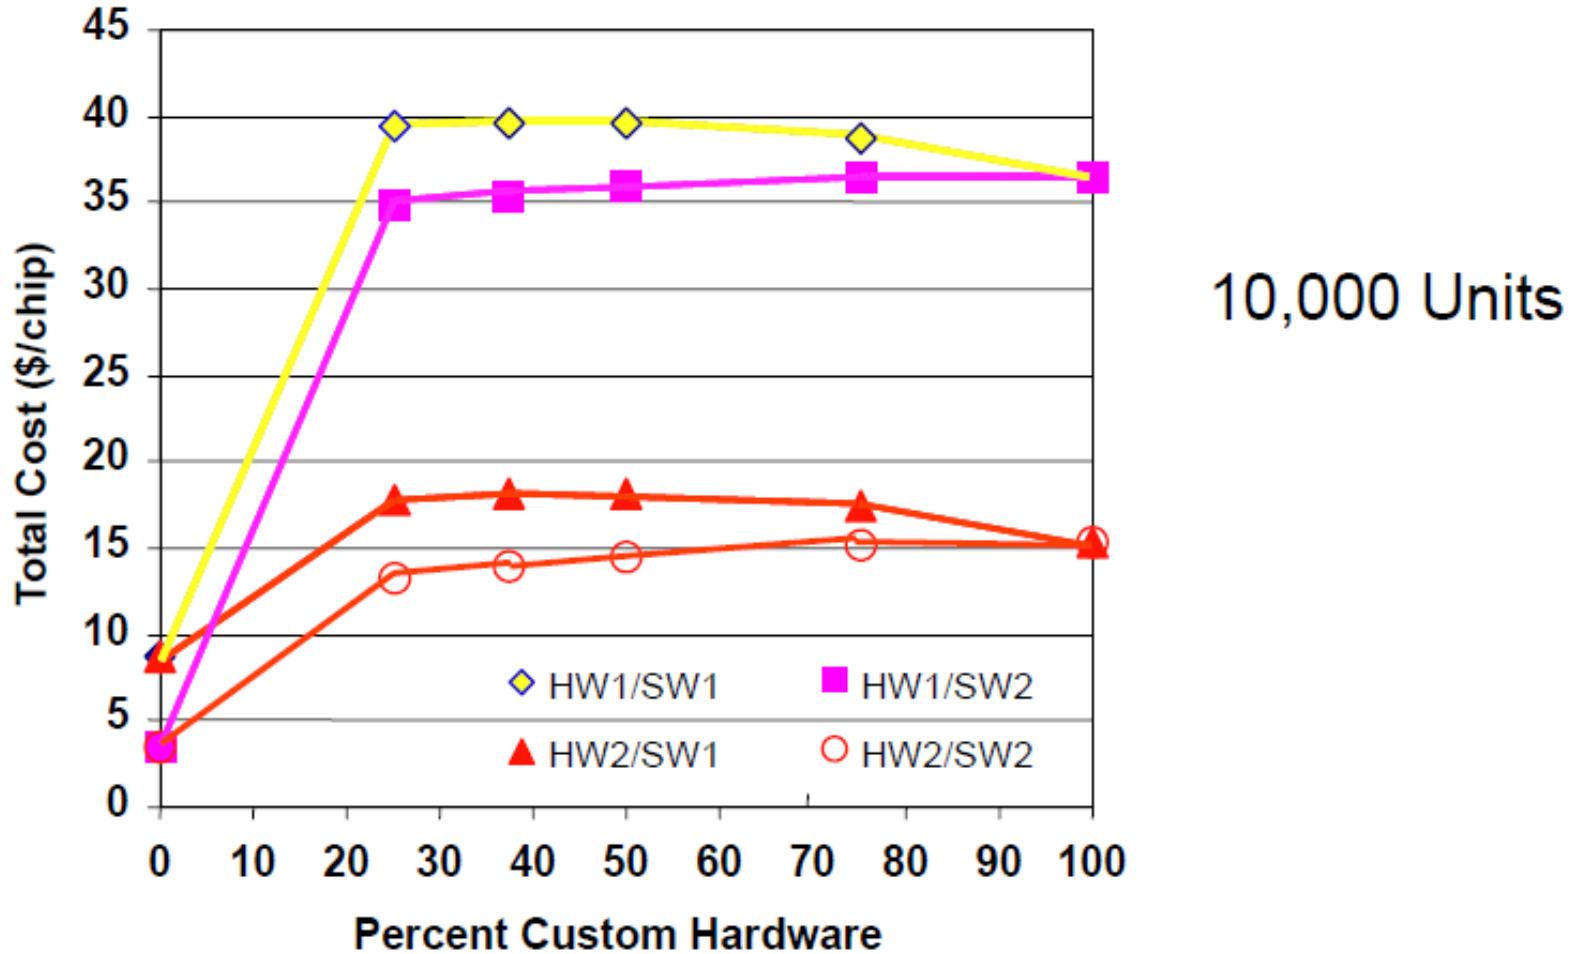
EXAMPLE: MIXED IMPLEMENTATION USING HW1 AND SW1



Reuse of:

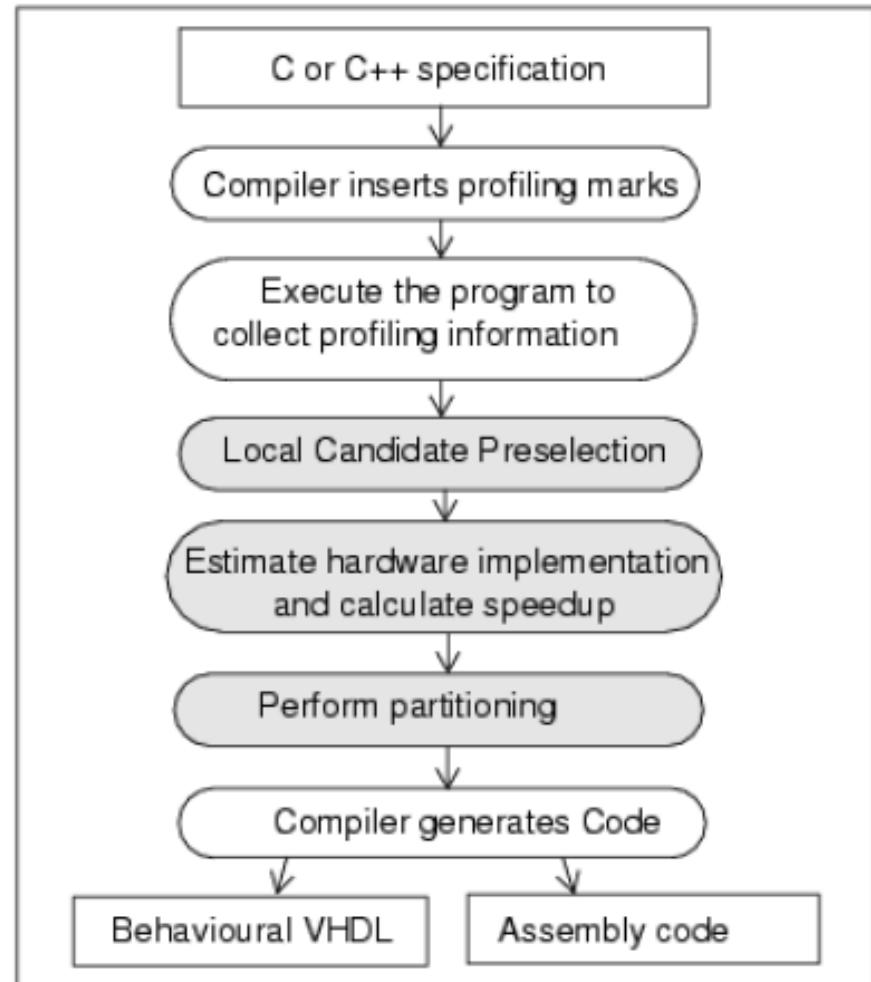
- Gate-level IP
- Code

EXAMPLE: TOTAL COST PER CHIP



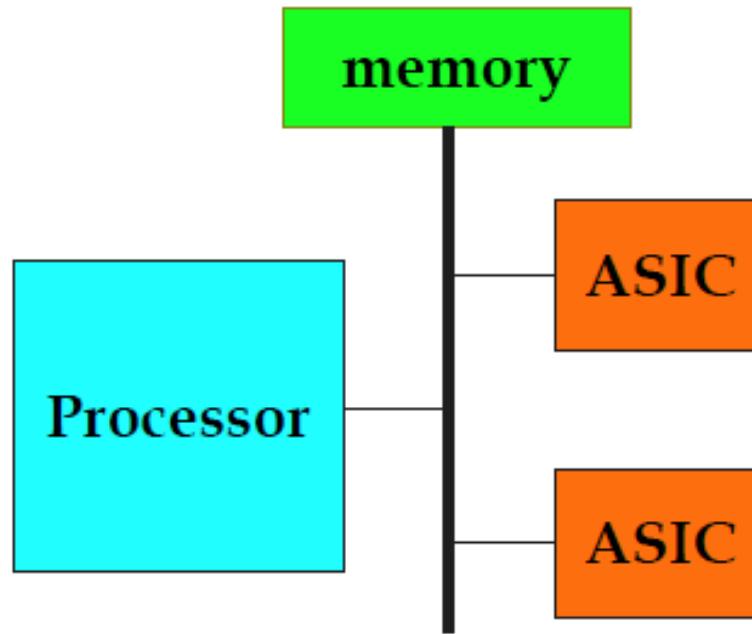
PARTITIONING ANALYSIS

- Result of compilation is synthesizable HDL and assembly code for the processor
- Compiler & profiler determine dependence and rough performance estimates



PARTITIONING STYLES

Simple architectural model: CPU + 1 or more ASICs on a bus



PARTITIONING STYLES

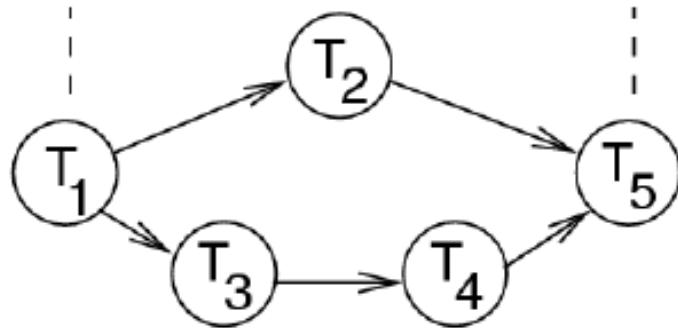
- **HW first approach**
 - start with all-ASIC solution which satisfies constraints
 - migrate functions to software to reduce cost
- **SW first approach**
 - start with all-software solution which does not satisfy constraints
 - migrate functions to hardware to meet constraints

SCHEDULING

25

SCHEDULING

- Goal: determine **when** and **where** a task is to be executed
- Inputs:
 - A task graph, node = task, edge = dependence



SCHEDULING

- Goal: determine **when** and **where** a task is to be executed
- Inputs:
 - A task graph, node = task, edge = dependence
 - Processor list, with different cost

- HW types H1, H2 and H3 with costs of 20, 25, and 30.
- Processors of type P.

SCHEDULING

- Goal: determine **when** and **where** a task is to be executed
- Inputs:
 - A task graph, node = task, edge = dependence
 - Processor list, with different cost
 - Task execution time on different processors

T	H1	H2	H3	P
1	20			100
2		20		100
3			12	10
4			12	10
5	20			100

SCHEDULING OUTPUTS

- For each task, determine
 - Its processor/HW assignment
 - Its starting time
- The overall hardware cost
 - The sum of costs of all needed units
- The overall schedule length
 - The finish time of the latest task
- The overall power consumption...

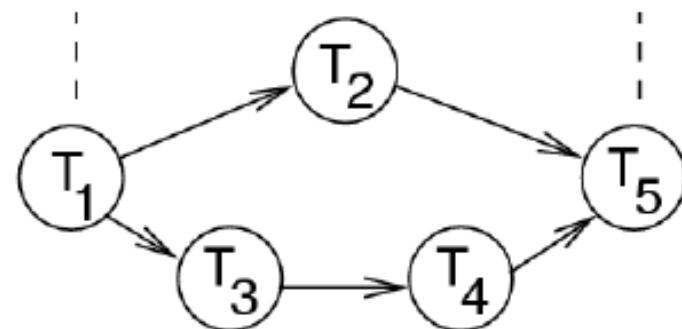
SCHEDULING CONSTRAINTS

- A task cannot start if any of its predecessor has not finished yet
- A processor/HW unit can execute only one task at any time

EXAMPLE 1:LOWEST-COST SCHEDULE

- To execute all 5 tasks we either
 - use P => cost = 50
 - use H1, H2, H3 => cost = 20+25+30 = 75
- Solution:
 - T1->P, start at time 0
 - T2->P, start at time 100
 - T3->P, start at time 200
 - T4->P, start at time 210
 - T5->P, start at time 220
 - Schedule length = 320

T	Cost	20	25	30	50
H1					P
1	20				100
2		20			100
3			12	10	
4			12	10	
5	20				100



EXAMPLE 2:SHORTEST SCHEDULE

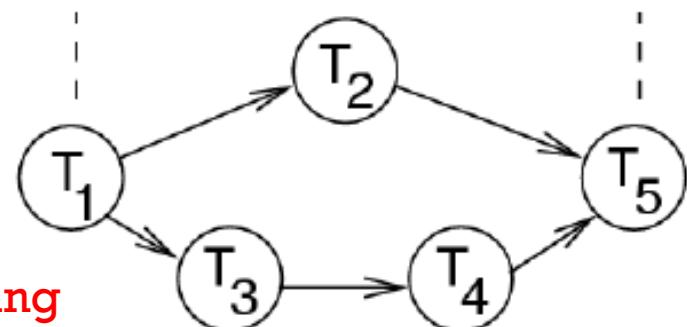
- Use the fastest HW for each task

- T1, T5->H1
- T2->H2
- T3, T4->P
- HW cost = $20+25+50 = 95$

- Solution:

- T1->H1, start at time 0
- T2->H2, start at time 20
- T3->P, start at time 20
- T4->P, start at time 30
- T5->H1, start at time 40
- Schedule length = 60

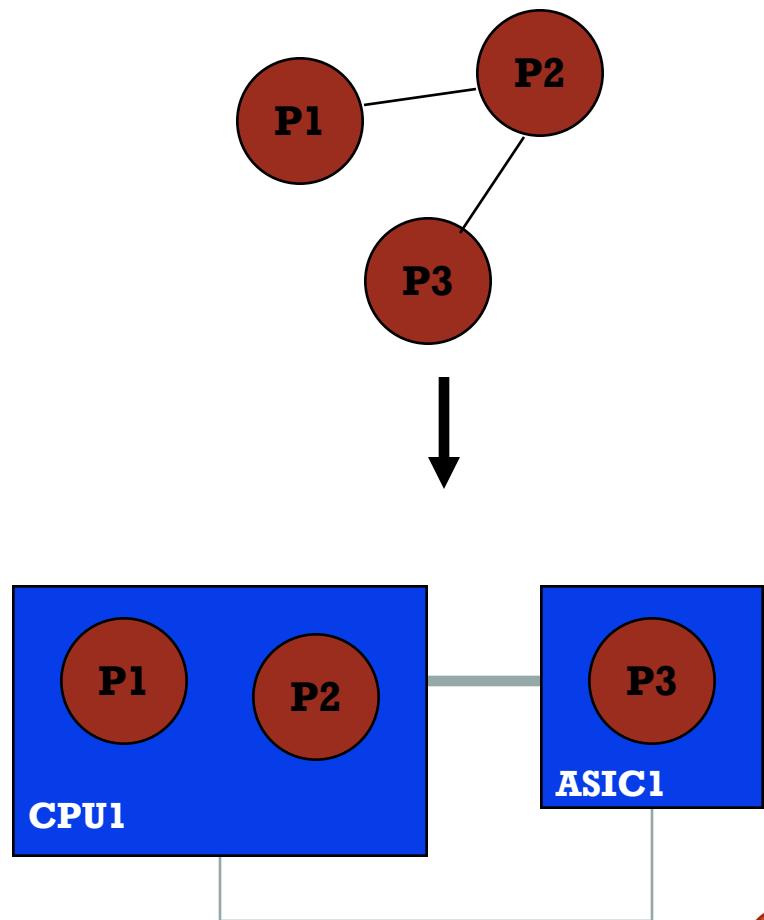
T	H1	H2	H3	P
1	20			100
2		20		100
3			12	10
4			12	10
5	20			100



The solution does not change even if assigning T2 on P takes a time of 10 instead of 100. Why?

COMMUNICATION COST

- Inter-processor communication cost is not zero!
- Must also schedule communication
- May impact schedule length significantly



CODESIGN MAPPING

Specification

Mapping

